Advanced Time Series Analysis: Computer Exercise 2

Anders Launer Bæk (s160159)

31 Oktober 2017

Sparring partners:

- Anja Liljedahl Christensen (s162876)

- Marie Mørk (s112770)

Part 1

There has been simulated $n = 3000$ samples of noise where $\epsilon_t \sim \mathcal{N}(0, 1)$. ϵ_t is used as noise input for simulations in part one.

The equation below shows the used parameters in the selected system: a SETAR(2,1,1) model. The parameters for the two regimes are defined by eqn. 1 and eqn. 2.

$$a_0 = [0.125, -0.125] \tag{1}$$

$$a_1 = [0.6, -0.4] \tag{2}$$

Simulation of the SETAR(2,1,1)

The SETAR(2,1,1) model is given by eqn. 3.

$$X_t = a_0^{(J_t)} + \sum_{i=1}^{k_{(J_t)}} a_i^{(J_t)} X_{t-i} + \epsilon^{(J_t)} \tag{3}$$

where J_t are the regime processes. The complete model are defined in eqn. 4.

$$X_t = \begin{cases} a_{0,1} + a_{1,1} X_{t-1} + \epsilon_t & for \quad X_{t-1} \leq 0 \\ a_{0,2} + a_{1,2} X_{t-1} + \epsilon_t & for \quad X_{t-1} > 0 \end{cases} \tag{4}$$

A simulation of the model (eqn. 4) is showed in fig. 1.

Estimate the parameters using conditional least squares

The following code snippet contain three functions: `Setar()`, `RSSSetar` and `PESetar` calculates the conditional means, the squared residuals and the total squared prediction error respectively.

```
# calculation of the conditional mean for the SETAR(2,1,1) model
Setar <- function(par, model) {
    # init conditional mean vector
    e_mean <- rep(NA, length(model))
    # loop through observations
    for (t in 2:length(model)) {
        if (model[t - 1] <= 0) {
            e_mean[t] <- par[1] + par[2] * model[t - 1]
```

Figure 1: One simulation of a SETAR(2,1,1) model.

```
        } else {
            e_mean[t] <- par[3] + par[4] * model[t - 1]
        }
    }
    # return the conditional mean vector
    return(e_mean)
}

RSSSetar <- function(par, model) {
    # conditional mean
    e_mean <- Setar(par, model)
    # calculate and return the squared residuals
    return((model - e_mean)^2)
}

# summed prediction error
PESetar <- function(par, model) {
    # conditional mean
    e_mean <- Setar(par, model)
    # calculate and return the objective function value
    return(sum((model - e_mean)^2, na.rm = TRUE))
}
```

The native R function `optim()` has been applied to estimate the parameters of the simulated process. The objective cost function is the total squared prediction error, which needs to be minimized.

eqn. 5 and eqn. 6 have been used as initial parameter input to optimization function. The "initial" conditional mean is computed with the initial parameters (eqn. 5 and eqn. 6) and illustrated in fig. 2 with the label: `M(x) initial`.

Table 1: Table shows the real parameters, the estimated parameters and their percentage diviation.

Parameter	Real value	Optimized value	Change in (%)
a0_1	0.125	0.1267946	1.4153889
a1_1	0.600	0.6127632	2.0828869
a0_2	-0.125	-0.1239285	-0.8645787
a1_2	-0.400	-0.4062514	1.5388065

$$a_0 = [0.1, -0.02] \tag{5}$$

$$a_1 = [0.4, -0.25] \tag{6}$$

The estimated parameters are listed in table 1.

The percentage deviation from the real parameters are within $\pm 2.0828869\%$. The optimized squared prediction error (30.0575719) is smaller than the real squared error (30.0607994) which is a subject of overfitting. Dividing data into train and test with proper cross validation techniques can solve the issue of overfitting.

Figure 2 illustrate the conditional means for the SETAR(2,1,1) model with initial parameters (eqn. 5 and eqn. 6), the real and estimated parameters (table 1).


Figure 2: Simulations of a SETAR(2,1,1) model with the initial, real and optimized parameters. The vertical lines in the transition of the regimes does not exist.

The initial parameters are not a good representation of the simulated SETAR model, as illustrated in fig. 2 with the label `M(x) initial`. But after estimating the parameters s.t. a minimization of the squared prediction error, the estimated parameters is a reasonable representation of the SETAR model.

Part 2

It has been chosen to create a grid based upon the two slope parameters, eqn. 2. The grid is based on a matrix with the size $(51, 51)$, where the centre element, the intersection of the diagonal and off-diagonal, is the optimal parameter value, and is highlighted with a blue dot. There is also illustrated a red dot, which shows the optimal parameters for the given subset of the simulation. The first axis is $a_{1,1}$ parameter and the second axis in the $a_{1,2}$ parameter.

The the deviation of the parameter values are $\pm 30\%$ (`max_change_p`) of the optimal value. The following function has been used to plot the contours for the different subsets: `plot_contours()`. General to the contour plots, close contour lines identicate a higher slope of the curvature of the objective function.

```r
# only change the slope par[2] and par [4]
plot_contours <- function(model, par = optimal_PE$par, change_p = max_change_p,
    nplot = resolution) {
    # create squence of nplot values with the optimized parameter(s) in center
    par_2_seq <- seq(par[2] - par[2] * change_p, par[2] + par[2] * change_p,
        len = nplot)
    par_4_seq <- seq(par[4] - par[4] * change_p, par[4] + par[4] * change_p,
        len = nplot)

    # create grid
    loess_melt <- expand.grid(par_2_seq, par_4_seq)

    # caculate values
    loess_melt$value <- sapply(1:nrow(loess_melt), function(i) {
        return(PESetar(par = c(par[1], loess_melt$Var1[i], par[3], loess_melt$Var2[i]),
            model = model))
    })

    # find optimal subset value
    min_val <- loess_melt[which.min(loess_melt$value), ]

    # return contour plot
    return(ggplot(loess_melt, aes(x = Var1, y = Var2, z = value)) + stat_contour(geom = "contour",
        alpha = 1/2, binwidth = 0.005) + geom_point(aes(x = par[2], y = par[4],
        color = paste0("optim,\nval=", round(PESetar(par = par, model = model),
            digits = 4))), alpha = 1/2) + geom_point(aes(x = min_val$Var1, y = min_val$Var2,
        color = paste0("best subset,\nval=", round(min_val$value, digits = 4))),
        alpha = 1/2) + labs(x = "a_1,1", y = "a_1,2", color = "parameter and value") +
        theme_TS())
}
```

The contour plots illustrates the objective function (total squared prediction error) with changes in each changeable parameter. The objective function is convex which means there is only one local/global minima.

N = 1:3000

Figure 3 shows the contour plot for the complete simulated model. The optimal parameters gives the global minima of the objective function, and the two dots are identical as illustrated.

The $a_{1,1}$ parameter will result in a higher objective function with equally changes in both parameter.

Figure 3: Contour plot the slope parameters for N = 1:3000.

N = 1:300

Figure 4 shows the contour plot with only $N = 1 : 300$ samples of the simulated model considered. The number of samples are decreased with a factor of 100 similar to the objective function.

The optimal parameters does not represent the minimum value of the of the objective function. There is a clear separation of the two dots and the their values are different, see labels in fig. 4.


Figure 4: Contour plot the slope parameters for N = 1:300.

N = 1:30

Figure 5 shows a contour plot of the first 30 samples. It is possible to see that the best set of parameter, in the given subset, is not within $\pm 2.0828869\%$ of the optimized parameters. In contrast to fig. 5 there red dot is placed on the boundary of the grid which indicates that the global minimum is not obtained.


Figure 5: Contour plot the slope parameters for N = 1:30.

N = 1001:1300

Figure 6 uses same number of simulated samples but in an different "location" of the simulated data. The optimal subset parameters is not within $\pm 30\%$ of the optimized parameters which is opposite to fig. 4.

N = 1001:1030

The best subset parameters in fig. 7 are not within $\pm 30\%$ of the optimized parameters.

Findings

- The optimization of the parameters must occur with respect to the wanted (sub)set of the simulated data.

- The biggest changes for the optimal subset parameters is for the $a_{1,2}$ parameter. In figures 4 and 5 the best subset $a_{1,2}$ parameter is smaller then the optimized value of $a_{1,2}$.
  Opposite in figures 6 and 7, the best subset $a_{1,2}$ parameter must be higher than the optimized $a_{1,2}$ parameter.

- The total squared prediction error scales well with the number of considered simulated data. This can identicate that the residuals are normally distributed over the entire simulated process.

Figure 6: Contour plot the slope parameters for N = 1001:1300.

Figure 7: Contour plot the slope parameters for N = 1001:1030.

Part 3

The chosen model for the doubly stochastic model, is an AR(2)-AR(4) model, see eqn. 7.

$$Y_t = \sum_{k=1}^{2} \left(\Phi_{t-(1-k)} Y_{t-k} \right) + \epsilon_t$$

$$\Phi_t - \mu = \sum_{n=1}^{4} \left(\phi_n \left(\Phi_{t-n} - \mu \right) \right) + \zeta_t \tag{7}$$

$$\Phi_t = \sum_{n=1}^{4} \left(\phi_n \left(\Phi_{t-n} - \mu \right) \right) + \zeta_t + \underbrace{\mu \left(1 - \sum_{n=1}^{4} (\phi_n) \right)}_{\delta_t}$$

Equation 8 shows eqn. 7 in a reparametrized state space format.

$$\begin{pmatrix} \Phi_t \\ \Phi_{t-1} \\ \Phi_{t-2} \\ \Phi_{t-3} \\ \delta_t \end{pmatrix} = \begin{pmatrix} \phi_1 & \phi_2 & \phi_3 & \phi_4 & 1 \\ 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} \Phi_{t-1} \\ \Phi_{t-2} \\ \Phi_{t-3} \\ \Phi_{t-4} \\ \delta_{t-1} \end{pmatrix} + \begin{pmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix} \delta_t$$

$$Y_t = \begin{pmatrix} Y_{t-1} & Y_{t-2} & 0 & 0 & 0 \end{pmatrix} \begin{pmatrix} \Phi_t \\ \Phi_{t-1} \\ \Phi_{t-2} \\ \Phi_{t-3} \\ \delta_t \end{pmatrix} + e_t \tag{8}$$

Simulate

The initial parameters for the simulation are given in eqn. 9.

$$\begin{aligned} n &= 500 \\ \mu &= 0.1 \\ \zeta &= 0.08 \\ \epsilon &= 0.4 \\ \delta_t &\sim \mathcal{N}(\mu, \zeta) \\ e_t &\sim \mathcal{N}(\mu, \epsilon) \end{aligned} \tag{9}$$

Figure 8 and figure 9 shows the process and the underlying process respectively.

Comment

- Figure 9 shows the underlying process of the doubly stochastic model and the its shows a rapidly increase after the first samples where the process more less stabilize around 0.2 afterwards.

- Choosing the initial parameters for this doubly stochastic model is critical and needs to be done with stability in mind.

- todo: Stability in the underlying process? and stability in the overlying process?

Figure 8: Simulated process of Y(t).

Figure 9: Simulated process of Phi(t).

Part 4

Following state space model is given, eqn. 10.

$$x_{t+1} = ax_t + v_t$$
$$y_t = x_t + e_t$$

(10)

where a is an unknown parameter and v_t and e_t are mutually uncorrelated white noise processes with their

variances σ_v^2 and σ_e^2.

Part 4a

The model from eqn. 10 is on state space form in eqn. 11

$$x_{t+1} = ax_t + v_t$$
$$y_t = x_t + e_t$$
(11)

Simulate

There has been simulated 20 time series with the initial parameters given in eqn. 12.

$$a = 0.4$$
$$\sigma_v^2 = 1$$
$$\sigma_e^2 = 1$$
$$v_t \sim \mathcal{N}(0, \sigma_v^2)$$
$$e_t \sim \mathcal{N}(0, \sigma_e^2)$$
(12)

Rewrite

Equation 11 has been reparametrized to state space form in eqn. 13.

$$\begin{pmatrix} x_{t+1} \\ a_{t+1} \end{pmatrix} = \begin{pmatrix} a_t & 0 \\ 0 & a_t \end{pmatrix} \begin{pmatrix} x_t \\ 1 \end{pmatrix} + \begin{pmatrix} v_t \\ 0 \end{pmatrix}$$
$$y_t = \begin{pmatrix} 1 & 0 \end{pmatrix} \begin{pmatrix} x_t \\ 1 \end{pmatrix} + e_t$$
(13)

a has been included in the state space equation, which makes it possible to estimate the parameter using the filter.

Part 4b

The following function (`ext_kalman()`) has been used to estimate a for the given sets of initial parameters.

```
##---------------------------------------------------------------
## EKF algorithm for use in Part 4 of computer exercise 2 in Advanced Time
## Series Analysis
##---------------------------------------------------------------

ext_kalman <- function(y, aInit = 0.5, aVarInit = 1, sigma.v = 1) {
    ## aInit : The starting guess of the AR coefficient estimate aVarInit : The
    ## initial variance for estimation of the AR coefficient sigma.v : Standard
    ## deviation of the system noise of x in the filter

    # Initialize---- Init the state vector estimate
    zt <- c(0, aInit)
    # Init the variance matrices
    Rv <- matrix(c(sigma.v^2, 0, 0, 0), ncol = 2)
    # sigma.e : Standard deviation of the measurement noise in the filter
```

```r
    Re <- 1

    # Init the P matrix, that is the estimate of the state variance
    Pt <- matrix(c(Re, 0, 0, aVarInit), nrow = 2, ncol = 2)
    # The state is [X a] so the differentiated observation function is
    Ht <- t(c(1, 0))
    # Init a vector for keeping the parameter a variance estimates
    aVar <- rep(NA, length(y))
    # and keeping the states
    Z <- matrix(NA, nrow = length(y), ncol = 2)
    Z[1, ] <- zt

    ## The Kalman filtering----
    for (t in 1:(length(y) - 1)) {
        # Derivatives (Jacobians)
        Ft <- matrix(c(zt[2], 0, zt[1], 1), ncol = 2)   # F_t-1
        # Ht does not change

        ## Prediction step
        zt = c(zt[2] * zt[1], zt[2])   #z_t/t-1 f(z_t-1/t-1)
        Pt = Ft %*% Pt %*% t(Ft) + Rv   #P_t/t-1

        ## Update step
        res = y[t] - zt[1]   # the residual at time t
        St = Ht %*% Pt %*% t(Ht) + Re   # innovation covariance
        Kt = Pt %*% t(Ht) %*% St^-1   # Kalman gain
        zt = zt + Kt * res   # z_t/t
        Pt = (diag(2) - Kt %*% Ht) %*% Pt   #P_t/t

        ## Keep the state estimate
        Z[t + 1, ] <- zt
        ## Keep the P[2,2], which is the variance of the estimate of a
        aVar[t + 1] <- Pt[2, 2]
    }
    return(list(zt = zt, Pt = Pt, Rv = Rv, aVar = aVar, Z = Z))
}
```

There has been executed a small test to find the number of needed samples for converges of the filter estimate. The test is based upon the worst case scenario where the initial variance of the filter and parameters is 10. The initial value of the parameter is $a = 0.5$.

Figure 10 shows the converges test for the filter estimate. The variance of the parameter and the estimated value of the parameter.

The filter is converged after the first 3000 samples, (figure 10) and will therefore only consider those first samples.

Table 2 and table 3 shows the initial variance values for each combination averages over the 20 simulations: The mean estimated value of a, the standard deviation of the estimated a, the mean value of the variance and the standard deviation of the variance.

a = 0.5

Table 2 shows the statistics for $a = 0.5$.

Figure 10: Converges test of the filter estimate.

Table 2: Statistic for filter estimates.

combination	sigma_v^2	sigma_a	a	a mean	a sd	a_var	a_var mean	a_var sd
1	10	1	NA	0.2224069	0.0155871	NA	0.0046646	6.06e-05
2	1	1	NA	0.4013185	0.0217119	NA	0.0002554	1.16e-05
3	10	10	NA	0.2212304	0.0156610	NA	0.0046843	6.11e-05
4	1	10	NA	0.4013258	0.0216101	NA	0.0002554	1.16e-05

a = -0.5

Table 3 shows the statistics for $a = -0.5$.

Table 3: Statistic for filter estimates.

combination	sigma_v^2	sigma_a	a	a mean	a sd	a_var	a_var mean	a_var sd
1	10	1	NA	0.2224069	0.0155871	NA	0.0046646	6.06e-05
2	1	1	NA	0.4013185	0.0217119	NA	0.0002554	1.16e-05
3	10	10	NA	0.2212304	0.0156610	NA	0.0046843	6.11e-05
4	1	10	NA	0.4013258	0.0216101	NA	0.0002554	1.16e-05

Comments

According to table 2 and table 3 there is not much of a difference in when selecting different initial parameter value of a. Therefore some common comments.

- The most important parameter, when using filter to estimate parameter(s), is the variance of the filter. The combinations with a low filter variance result in a reasonable parameter estimate, close to the original value of a, regarding the initial value of variance of the parameter.

- In both combinations where $\sigma_v^2 = 1$ is the estimated value of a slightly above the original value of a.

Improvements

Estimate a parameter with an extended Kalman filter will introduce a bias estimate coursed by the properties of the the filter.